

The Sagan Information Scale

Knowledge as the Fifth Lens

Volume 5 in the Civilization Scales Series:

Thirteen Books on Twelve Classification Frameworks + Synthesis

by James Edmund Carpenter II with xAI & Kimi-AI

Table of Contents

Introduction: The Knowledge Lens

Chapter 1: The Cosmic Messenger

Chapter 2: The Cosmic Connection - The Original Formulation

Chapter 3: The Alphabet of Knowledge - How the Scale Works

Chapter 4: Where Humanity Stands - From H to S in Fifty Years

Chapter 5: The Thermodynamics of Information - From Maxwell's
Demon to Landauer's Limit

Chapter 6: The Ultimate Laptop - Seth Lloyd and the Physical
Limits of Computation


Chapter 7: Dyson's Eternal Intelligence - Information Preservation
Across Cosmic Time

Chapter 8: Criticisms and Limitations - Does Information Equal
Knowledge?

Chapter 9: The Knowledge Lens - Preparing for the Epilogue
Twelve-Scale Synthesis

Notes on Sources and Further Reading

Rolling Notation Guide (Series-Wide)



Introduction: The Knowledge Lens

In the spring of 1973, a thirty-eight-year-old astronomer named Carl Sagan published a book that would transform the public understanding of extraterrestrial life and, almost as an afterthought, introduce a second dimension to the classification of civilizations. The book was *The Cosmic Connection: An Extraterrestrial Perspective*, and its author was already one of the most visible scientists in America — a regular on *The Tonight Show Starring Johnny Carson*, a consultant to NASA, and the public face of the Search for Extraterrestrial Intelligence. *The Cosmic Connection* would become a bestseller, establishing Sagan as the preeminent science communicator of his generation.

But buried within its pages, alongside discussions of Mars, Venus, and the search for radio signals from distant stars, was something else: a proposal that the amount of information a civilization commands might be as important as the amount of energy it consumes. Sagan called this the Information Mastery Scale, and it was the first serious attempt to create a civilization classification framework based not on watts or parsecs but on bits — the fundamental unit of information.

The idea was both simple and profound. Sagan assigned each letter of the alphabet to represent an order of magnitude in information content: Type A civilizations possessed 10^6 unique bits, Type B civilizations 10^7 bits, and so on up to Type Z at 10^{31} bits. Earth

in 1973, Sagan estimated, was a Type H civilization — approximately 10^{13} bits of information, drawn largely from the holdings of the world's great libraries. But Sagan went further. He proposed a two-dimensional classification that combined energy (the Kardashev Scale) with information (his letter scale), producing composite ratings like “0.7 H” for 1973 Earth and “1.5 J” or “1.8 K” for the civilizations humanity might first contact. This is the Sagan Information Scale, and it is the subject of the book you are now reading. It is the fifth lens in our growing toolkit — after Kardashev's energy lens (Book 1), Zubrin's spatial lens (Book 2), Clarke's tool lens (Book 3), and Barrow's inward lens (Book 4). Where the previous four scales measured what civilizations *can do* — how much energy they command, how far they can travel, what tools they can build, how small they can manipulate — the Sagan scale measures what civilizations *know*. And in an age of artificial intelligence, big data, and the exponential growth of information, this may be the most relevant of all the scales we will examine.

The Sagan scale connects directly to some of the deepest ideas in modern physics: the thermodynamics of information, the resolution of Maxwell's demon, the Landauer limit on the energy cost of computation, and the ultimate physical limits on how much information can be stored and processed. It also connects to one of the most provocative ideas in cosmological eschatology: Freeman

Dyson's concept of eternal intelligence — the possibility that advanced civilizations can survive indefinitely, not by commanding ever more energy, but by learning to process information with ever greater efficiency as the universe cools toward absolute zero.

By the end of this book, we will have examined all of these connections in depth. We will have traced Sagan's scale from its origins in a 1973 bestseller, through the physics of information that gives it theoretical foundation, to its implications for understanding the long-term fate of civilizations in an expanding universe. We will have seen why information may be, as Sagan suggested, "more significant" than energy as a criterion for classifying civilizations. And we will have added a fifth lens to our growing toolkit — a way of seeing civilizations not merely as engines of power, but as engines of knowledge.

Chapter 1

The Cosmic Messenger

Carl Edward Sagan was born on November 9, 1934, in Brooklyn, New York, the son of Samuel Sagan, a Russian-Jewish immigrant who worked in the garment industry, and Rachel Gruber, a homemaker. The family lived in a modest apartment in the Bensonhurst neighborhood, and young Carl's path to science was paved not by privilege but by an insatiable curiosity about the natural world and a mother who encouraged his questions rather than brushing them aside.

As a child, Sagan was fascinated by the stars. He would attend the 1939 New York World's Fair and lie on the grass looking up at the night sky, wondering what those points of light were and whether anyone was looking back. This childhood wonder would never leave him. It would drive him to a career that combined rigorous scientific research with passionate public communication, making him arguably the most effective science popularizer since Galileo. Sagan attended the University of Chicago on a full scholarship, where he earned a Bachelor of Science degree in 1955 and a Master of Science in 1956, both in physics. He completed his PhD in astronomy and astrophysics in 1960, working under the supervision of Gerard Kuiper, one of the giants of twentieth-century planetary science. Kuiper was a Dutch-American astronomer who had made fundamental contributions to the study

of the solar system, including the discovery of Miranda, a moon of Uranus, and Nereid, a moon of Neptune. He was also one of the first scientists to propose the existence of a belt of cometary bodies beyond the orbit of Neptune — what is now known as the Kuiper Belt. Sagan’s doctoral thesis, “Physical Studies of the Planets,” laid the groundwork for his lifelong interest in planetary atmospheres, surfaces, and the potential for extraterrestrial biology. Sagan’s early research was extraordinarily productive and spanned multiple fields. In 1962, he demonstrated that Venus’s high surface temperature was caused by a runaway greenhouse effect — a finding that has profound implications for understanding climate change on Earth. Venus and Earth are approximately the same size and composition, yet Venus has a surface temperature of 462 degrees Celsius — hot enough to melt lead — while Earth maintains a habitable average of 15 degrees Celsius. The difference, Sagan showed, is the atmosphere. Venus’s thick carbon dioxide atmosphere traps heat through the greenhouse effect, while Earth’s thinner atmosphere, balanced by biological processes like photosynthesis, maintains a temperature suitable for liquid water and life. This insight — that a planet’s atmosphere is a delicate, dynamic system that can be pushed into radically different states — would later inform Sagan’s thinking about nuclear winter, climate change, and the fragility of civilizations.

Sagan also showed that the seasonal darkening on Mars was likely caused by windblown dust, not vegetation as many astronomers had hoped. He proposed that the building blocks of life could be synthesized from simple chemicals under conditions mimicking the early Earth's atmosphere — extending the famous Miller-Urey experiment, which had demonstrated the synthesis of amino acids from methane, ammonia, hydrogen, and water vapor subjected to electrical discharges. Sagan's variations included ultraviolet radiation and different atmospheric compositions, showing that the origin of life's building blocks might have been a robust process occurring under a wide range of conditions.

Sagan was among the first scientists to take seriously the possibility of extraterrestrial life, at a time when the topic was considered fringe at best by the scientific establishment. In 1966, he co-authored *Intelligent Life in the Universe* with the Soviet astrophysicist Iosif Shklovsky — a comprehensive survey of the scientific basis for SETI that became the foundational textbook of the field. The book covered everything from the origin of life on Earth to the search for radio signals, the possibility of interstellar travel, and the philosophical implications of contact. It was translated into multiple languages and influenced a generation of astronomers, biologists, and engineers who would go on to make SETI a respected scientific discipline.

But Sagan's greatest impact came not from his research papers but from his ability to communicate science to the public. In 1973, he published *The Cosmic Connection*, the book that introduced both his decimal extension of the Kardashev Scale and the Information Mastery alternative. In 1978, he won the Pulitzer Prize for general non-fiction for *The Dragons of Eden: Speculations on the Evolution of Human Intelligence*, a book that explored the evolution of the human brain and the nature of consciousness through the lens of comparative neuroanatomy and evolutionary biology. The book was praised for its ability to make complex neuroscience accessible to a general audience while maintaining scientific rigor.

In 1980, Sagan co-wrote and narrated the television series *Cosmos: A Personal Voyage*, produced by PBS affiliate KCET Los Angeles and broadcast in more than sixty countries to an audience of over half a billion people. It remains the most-watched series in the history of American public television. The thirteen-episode series covered the origin of life, the evolution of the cosmos, the nature of consciousness, the search for extraterrestrial intelligence, and the responsibility of humanity as caretakers of Earth. Its signature phrase — “We are a way for the cosmos to know itself” — has become one of the most quoted lines in popular science. The accompanying book, *Cosmos*, spent more than seventy weeks on

the *New York Times* bestseller list and has been translated into dozens of languages.

Sagan was not merely a communicator. He was a participant in some of the most important scientific projects of his era. He designed the Pioneer plaque — the first physical message from humanity to extraterrestrial intelligence, affixed to the Pioneer 10 and 11 spacecraft. The plaque, designed with the artist Linda Salzman Sagan (his wife at the time), shows a naked human male and female, a diagram of the hydrogen atom, a map of the Sun's position relative to fourteen nearby pulsars, and a silhouette of the spacecraft itself. The message was intended to be comprehensible to any civilization with sufficient scientific knowledge, regardless of their language or culture. The choice of pulsars as a location marker was particularly clever: pulsars have unique rotation periods, so an advanced civilization could identify the Sun's position by matching the depicted periods to observed pulsars. Sagan chaired the committee that created the Voyager Golden Record — a gold-plated copper phonograph record carrying sounds and images of Earth, including greetings in fifty-five languages, music from diverse cultures (ranging from Bach to Chuck Berry), natural sounds like thunder and bird songs, and 116 photographs of life on our planet. The record is now traveling beyond the solar system on the Voyager 1 and 2 probes, having crossed the heliopause into interstellar space in 2012 and 2018

respectively. Voyager 1 is now more than 24 billion kilometers from Earth — the most distant human-made object in existence — and the Golden Record may survive for billions of years in the cold vacuum of interstellar space, long after the Sun has become a red giant and consumed the inner planets.

Sagan advised NASA on the Viking missions to Mars — the first spacecraft to successfully land on the Red Planet and return data — as a member of the imaging science team. He was a tireless advocate for SETI, arguing that the search for extraterrestrial intelligence was not science fiction but a legitimate and urgent scientific endeavor. He was an early and vocal proponent of action on climate change and nuclear disarmament — causes he pursued with the same passion he brought to astronomy. His 1983 paper on nuclear winter, co-authored with Richard Turco, Brian Toon, Thomas Ackerman, and James Pollack, demonstrated that a full-scale nuclear exchange would inject enough soot and dust into the stratosphere to block sunlight and cause a catastrophic drop in global temperatures — potentially ending civilization regardless of the direct effects of the explosions.

The Sagan Information Scale emerged from this unique combination of scientific rigor and cosmic imagination. Sagan was not merely proposing a classification system. He was asking a profound question: what is the most important thing that advanced civilizations have in common? Kardashev had answered: energy.

Sagan answered: information. Both answers are partial. Both are insightful. And both are essential to the multi-dimensional framework we are building in this series.

Chapter 2

The Cosmic Connection

The Original Formulation

The Cosmic Connection was published in 1973 by Anchor Press/Doubleday. The book emerged from a series of lectures and essays Sagan had written over the preceding decade, and it synthesized his thinking about extraterrestrial life, interstellar communication, and the nature of intelligence into a single accessible volume. The book's central argument was that the search for extraterrestrial intelligence was not merely a scientific curiosity but a civilizational imperative — that understanding our place in the cosmos required us to search for others who might share it.

The book was written at a pivotal moment in the history of SETI. NASA was planning the Viking missions to Mars, scheduled for launch in 1975, and Sagan had been deeply involved in their design as a member of the imaging science team. He had argued for including biology experiments on the landers — the Labeled Release, Gas Exchange, and Pyrolytic Release experiments that would search for metabolic activity in Martian soil. The first systematic radio telescope searches for extraterrestrial signals — Project Ozma in 1960, led by Frank Drake at the Green Bank Observatory in West Virginia, and its successors — had yielded no

confirmed detections, but the technology was improving rapidly. Drake's original search had targeted two nearby stars, Tau Ceti and Epsilon Eridani, listening at the 21-centimeter hydrogen line frequency — a natural “watering hole” where radio astronomers and alien civilizations alike might gather to communicate. The discovery of pulsars in 1967 by Jocelyn Bell Burnell, then a graduate student at the University of Cambridge, and her thesis advisor Antony Hewish, had briefly raised hopes of an artificial origin. The signals were so regular — pulsing with a precision of better than one part in ten million — that they were initially dubbed “LGM-1” for “Little Green Men” as a half-joking reference to their possible artificial origin. Bell Burnell and Hewish soon identified the signals as naturally occurring emissions from rapidly rotating neutron stars — incredibly dense stellar remnants that emit beams of radiation like cosmic lighthouses. The discovery earned Hewish the 1974 Nobel Prize in Physics (a controversial decision, as Bell Burnell's crucial role was not recognized), but it also taught SETI researchers an important lesson: nature can produce signals that mimic artificial sources, and extraordinary claims require extraordinary evidence. The development of large radio telescopes was making it possible to search for signals from increasingly distant stars. The Arecibo Observatory in Puerto Rico, completed in 1963, featured a 305-meter dish — the largest single-aperture telescope in the world

until the construction of China's FAST telescope in 2016. The Green Bank Telescope in West Virginia, completed in 2000, had a 100-meter by 110-meter offset dish that made it ideal for SETI observations. And the development of digital signal processing technology was making it possible to search millions of frequency channels simultaneously — a vast improvement over the single-channel receivers used in Project Ozma.

The time was ripe for a book that would bring SETI to a broad audience, and Sagan was the person to write it. *The Cosmic Connection* covers a wide range of topics: the possibility of microbial life on Mars, the chemistry of the origin of life, the design of interstellar messages and the problem of making them comprehensible to alien minds, the psychology of belief in extraterrestrials and the history of UFO sightings, the political and ethical implications of contact with a more advanced civilization, and the possibility that intelligent life might be far more common — or far more rare — than we imagine.

But for our purposes, the most important chapters are those in which Sagan discusses the classification of civilizations. Before proposing his information scale, Sagan addressed a limitation of the Kardashev Scale that he found genuinely problematic: the gaps between the three types were so vast that the scale lacked the resolution to describe civilizations in transition. “The energy gap between a Type I and a Type II civilization, or between a Type II

and a Type III civilization is enormous,” Sagan wrote. “It seems useful, if the matter is to be considered seriously, to have a finer degree of discrimination.”

Sagan’s solution was to introduce decimal subdivisions using Arabic numerals. He suggested that Type 1.0 would correspond to a civilization using 10^{16} watts, Type 1.1 to 10^{17} watts, and so on, with each increment of 0.1 representing a factor of ten increase in power. This interpolation produces the formula $K = (\log_{10}(P) - 6) / 10$, where K is the Kardashev rating and P is the power consumption in watts.

As Robert H. Gray documented in his meticulous 2020 paper in *The Astronomical Journal*, Sagan never wrote this formula explicitly in *The Cosmic Connection*. What appears in the book is a verbal description of decimal subdivisions that is mathematically consistent with the equation but does not state it directly. Gray traced the formula’s emergence to later writers, possibly crystallizing from online discussions in the early 2000s — perhaps in a 2004 Wikipedia article, though no single originator has been identified. Gray concluded that while the formula is universally attributed to Sagan, its precise authorship remains unknown — it appears to be a communal formalization of Sagan’s verbal idea, refined by generations of SETI researchers who found the compact notation useful.

Sagan estimated that 1973 Earth, consuming approximately 10 terawatts (10^{13} watts), rated about $K = 0.7$ on this scale. This figure has become one of the most widely cited numbers in discussions of the Kardashev Scale. Gray's bibliographic investigation confirmed that Sagan's original estimate was somewhat rough — Sagan did not provide detailed calculations — and subsequent analyses have refined it. In 2021, global energy consumption was approximately 595 exajoules per year, equivalent to about 18.87 terawatts of continuous power, yielding a Kardashev rating of approximately 0.73. In 2024, with energy consumption reaching approximately 620 exajoules per year, the rating had crept to approximately 0.74 — a modest increase over half a century.

Having refined the energy scale, Sagan then proposed the complementary dimension that is the subject of this book. “But there may be more significant ways to characterize civilizations than by the energy they use,” he wrote. “An important criterion of a civilization is the total amount of information that it stores. This information can be in the form of the contents of libraries, or of the total hereditary information contained in the nucleotide sequences of the DNA of the organisms comprising the civilization.”

Sagan assigned the letter A to represent 10^6 unique bits of information — “less than any recorded human culture,” he noted, “and a good beginning point.” Each successive letter represented

an order of magnitude increase: $B = 10^7$, $C = 10^8$, and so on, up to $Z = 10^{31}$ bits. The choice of 10^6 as the baseline was deliberate: it corresponds to a civilization barely above the threshold of recorded culture, providing a floor against which all others can be measured.

To estimate Earth's 1973 rating, Sagan turned to libraries. The world's largest libraries at the time — the Library of Congress (approximately 71 million items), the British Museum, the Bibliothèque Nationale — held collections on the order of 10^7 to 10^8 items including books, manuscripts, maps, and recordings. Sagan estimated the total information content of these holdings at approximately 10^{13} to 10^{15} bits, placing 1973 Earth at Type H (10^{13} bits) or possibly Type I (10^{14} bits). In his composite classification, he described Earth as “0.7 H” — a civilization with Kardashev rating 0.7 and information rating H.

The composite notation is one of Sagan's most elegant innovations. It captures the two-dimensional nature of civilizational development in a compact form. A “0.7 H” civilization commands modest energy and modest information. A “1.5 J” civilization commands planetary-scale energy and substantially more information. A “3 Q” civilization — Sagan's estimate for a galactic supercivilization of a million worlds, each with a thousand times Earth's information — commands galactic-scale energy and information at the 10^{26} bit level. And a “4 Z”

civilization, at the upper limit of the scale, would command universal-scale energy and 10^{31} bits of information. Crucially, Sagan emphasized that the energy and information axes are “not strictly interdependent.” A civilization could rate high on one scale and low on the other. A civilization that had achieved extraordinary computational capabilities at relatively modest energy levels — through exceptional efficiency, or through technologies that extract maximum information from minimal power — would rate high on the information scale without necessarily rating high on the energy scale. This insight — that information and energy are complementary rather than equivalent metrics — is the foundation of everything that follows in this book, and it directly connects to the Efficiency Thesis introduced in Book 1 and developed further in Book 8 (STEM Compression Scale).

Chapter 3

The Alphabet of Knowledge

How the Scale Works

Sagan's information scale assigns each letter of the alphabet to a specific quantity of unique bits, with each step representing a factor of ten increase. The full scale runs from A to Z, covering twenty-six orders of magnitude from 10^6 to 10^{31} bits. This range was not chosen arbitrarily. The lower bound of 10^6 bits represents a civilization at the very threshold of recorded culture — just barely above the information content of individual human memory. The upper bound of 10^{31} bits represents a quantity of information so vast that it exceeds what Sagan believed any civilization could practically store or utilize.

Type A represents 10^6 bits — less than any recorded human culture. This would correspond to a very small tribal society with minimal technology and no written language, perhaps a group of hunter-gatherers numbering fewer than a hundred individuals. At this level, all knowledge is held in individual human memories and transmitted through oral tradition. The total information content is limited by the number of individuals and the capacity of human memory — which, as Thomas K. Landauer estimated in his 1986 study, is approximately 10^9 bits at midlife for a well-educated individual. A Type A civilization has not yet developed the

external memory systems — writing, counting, measurement — that allow knowledge to accumulate across generations.

Type B represents 10^7 bits — a small human society with basic oral traditions and simple tools. This might correspond to a tribal society with developed oral histories, totemic systems, and ceremonial knowledge passed through generations. Many indigenous cultures, past and present, operated at this level for millennia, accumulating vast bodies of knowledge about their environments — plant species, animal behavior, weather patterns, medicinal properties — all encoded in songs, stories, and rituals rather than in writing.

Type C represents 10^8 bits — a society with emerging written language and organized record-keeping. The earliest writing systems — Sumerian cuneiform in Mesopotamia around 3400 BCE, Egyptian hieroglyphs around 3200 BCE, Chinese oracle bone script around 1200 BCE, and the Indus script (still undeciphered) around 2600 BCE — represented humanity's first transition from purely oral transmission to externalized memory. Writing allowed administrators to track trade, taxes, and labor; priests to record religious doctrines; and scholars to accumulate knowledge that could outlive any individual. The total information content of these early literate civilizations was modest by modern standards — a few thousand clay tablets, a few hundred papyrus

scrolls — but it represented a qualitative leap in humanity's ability to preserve and transmit knowledge.

Type D represents 10^9 bits — a civilization with a substantial written tradition, comparable to Classical Greece at its height in the fifth and fourth centuries BCE. The physicist Philip Morrison of MIT estimated the information content of Classical Greek civilization at approximately this level, taking into account the epic poems of Homer (the *Iliad* and the *Odyssey*), the dialogues of Plato, the histories of Herodotus and Thucydides, the plays of Sophocles, Euripides, and Aristophanes, the philosophical works of Aristotle, the mathematical treatises of Euclid and Archimedes, and the medical writings of Hippocrates. Much of Greek knowledge was lost in the centuries following the Roman conquest — the Library of Alexandria was destroyed, manuscripts were burned or left to decay — but enough survived, primarily through Arab scholars who preserved and extended Greek learning during Europe's Dark Ages, to form the foundation of Western intellectual tradition.

Type E represents 10^{10} bits — a major ancient civilization with extensive libraries and administrative systems. Imperial Rome, Han China, and the Mauryan Empire in India all operated at approximately this level, maintaining vast bureaucracies that generated and preserved enormous quantities of written material. The Roman Empire's administrative records, legal codes, military

dispatches, and literary works collectively represented a significant accumulation of information, though much of it was lost after the empire's collapse.

Type F represents 10^{11} bits — a civilization approaching the information density of medieval Europe or imperial China at their peaks. Medieval European monasteries were the primary repositories of knowledge during this period, with scriptoria producing handwritten manuscripts at a rate of perhaps a few thousand per century across the entire continent. Imperial China's civil service examination system, which required candidates to memorize vast quantities of classical texts, created an enormous incentive for the preservation and transmission of written knowledge. The Song Dynasty (960-1279 CE) saw major advances in printing, cartography, and scientific observation, accumulating information at a rate that may have approached the lower end of this range.

Type G represents 10^{12} bits — the threshold of the print era. Johannes Gutenberg's invention of movable-type printing around 1440 CE initiated a slow but steady acceleration in the rate of information accumulation. Before Gutenberg, a skilled scribe might produce a single book in several months of labor. After Gutenberg, a printing press could produce hundreds of copies in the same time. By 1500, approximately twenty million books had been printed in Europe — more than all the manuscripts produced

by European scribes in the previous millennium. The information content of civilization had crossed a threshold.

Type H represents 10^{13} bits — Earth in 1973, by Sagan's estimate. This corresponds to the information holdings of the world's major libraries at that time — approximately 10^8 published works, each containing roughly 10^5 bits of unique information, yielding a total of approximately 10^{13} bits. Sagan acknowledged that this was a rough estimate and that different counting methods would yield different results.

Type I represents 10^{14} bits — the information frontier of the early digital era, when computers began to supplement paper-based storage in significant ways. By the mid-1980s, digital databases, electronic journals, and computerized records were adding measurably to humanity's total information store.

Type J represents 10^{15} bits — a civilization with extensive digital infrastructure, comparable to Earth in the late 1990s. The World Wide Web, invented by Tim Berners-Lee in 1989 at CERN and widely adopted by the mid-1990s, transformed the rate at which information could be created, shared, and stored. By 2000, the Web contained billions of pages and was growing exponentially.

Type K represents 10^{16} bits — a civilization with mature internet and global digital connectivity, comparable to Earth in the 2010s. Search engines like Google (founded 1998), social media platforms like Facebook (2004) and Twitter (2006), cloud

computing services like Amazon Web Services (2006), and mobile devices had by this point created an information ecosystem of unprecedented scale and complexity.

Type L represents 10^{17} bits — a civilization with massive data generation capabilities, entering the big data era. Machine learning, artificial intelligence, and the Internet of Things began producing data at rates that dwarfed human-created content.

Type M represents 10^{18} bits — a civilization generating exabytes of unique information annually. Michael Lesk's 1997 study estimated that the total human knowledge base had grown to approximately this level by the late twentieth century.

Type N represents 10^{19} bits — a civilization at the zettabyte threshold of data generation.

Type O represents 10^{20} bits — a civilization with infrastructure for extensive space colonization, in speculative extensions of the scale.

Type P represents 10^{21} bits — a civilization capable of stellar engineering.

Type Q represents 10^{22} bits — a galactic-scale civilization, in Sagan's conception. He estimated that a galactic supercivilization comprising a million worlds, each with a thousand times Earth's information, would rate "3 Q" — combining Kardashev Type III energy with Type Q information.

Type R represents 10^{23} bits

Type S represents 10^{24} bits. — Earth in 2024-2025, based on current data generation rates.

Type T represents 10^{25} bits.

Type U represents 10^{26} bits.

Type V represents 10^{27} bits.

Type W represents 10^{28} bits.

Type X represents 10^{29} bits.

Type Y represents 10^{30} bits.

Type Z represents 10^{31} bits — the ultimate category. Sagan believed no civilization had reached this level. He speculated that so much unique information would exceed that of all intelligent species in a galactic supercluster, and noted that the universe is not old enough for information to be effectively exchanged over such large distances.

The composite classification system combines the Kardashev energy rating with the information letter rating in the format “K_I”. Sagan estimated that first contact would involve civilizations rated “1.5 J” or “1.8 K” — advanced enough to communicate across interstellar distances, but not so advanced as to be incomprehensible. The composite system also reveals possibilities for “unbalanced” civilizations — a “0.5 Q” civilization with galactic-scale information but sub-planetary energy, or a “2.0 D” civilization with stellar-scale energy but limited information.

Chapter 4

Where Humanity Stands

From H to S in Fifty Years

When Sagan published *The Cosmic Connection* in 1973, he estimated Earth's information content at approximately 10^{13} to 10^{15} bits, placing humanity at Type H or possibly Type I on his scale. This estimate was based on the holdings of the world's great libraries — the Library of Congress (founded 1800, approximately 71 million items at the time), the British Library (founded 1973 from the collections of the British Museum, approximately 13 million items), the Bibliotheque Nationale de France (founded 1368, approximately 14 million books and periodicals), and their counterparts around the world.

Even at the time, other scholars produced different estimates using different methodologies. Thomas Kardashev, in his 1964 paper, had estimated 10^{14} bits for global information holdings, counting approximately 10^8 publications at 10^6 bits each. This method — counting the number of published works and multiplying by an average information content per work — was straightforward but prone to error, as different types of publications contain vastly different amounts of information. A scientific monograph might contain 10^7 bits of dense, novel information, while a popular magazine might contain 10^5 bits of largely redundant content.

Philip Morrison of MIT, who was Sagan's colleague and collaborator on SETI research, estimated the information content of Classical Greek civilization at approximately 10^9 bits — one ten-thousandth of Sagan's estimate for 1973 Earth. This estimate was based on counting the surviving works of Greek literature, philosophy, and science and extrapolating to account for lost works. Morrison's estimate illustrates how rapidly information can accumulate: the entirety of Classical Greek civilization, which laid the foundations of Western philosophy, mathematics, science, and democracy, contained less information than a single modern university library.

Thomas K. Landauer, a cognitive scientist at Bell Communications Research, estimated the “functional information content of human memory” at 10^9 bits at midlife — a figure that represents the total knowledge an educated individual can actively recall and use. This is not the total information stored in the brain (which is much larger, including sensory memories, motor skills, and unconscious knowledge) but the subset that is available for conscious reasoning and communication. Landauer estimated the “total information content of human memory” at 10^{10} bits, taking into account the storage capacity of neural connections. A single human brain, in other words, contains approximately as much information as all of Classical Greek civilization — a remarkable fact that underscores

both the sophistication of biological information processing and the vastness of humanity's collective knowledge.

What Sagan could not have anticipated was the speed at which humanity's information store would grow in the decades following 1973. In the fifty years since *The Cosmic Connection* was published, the growth of information has been not merely exponential but super-exponential — driven by the internet, personal computing, mobile devices, cloud storage, social media, streaming video, and the Internet of Things.

The numbers are staggering. In 2010, the total amount of data created, captured, copied, and consumed globally was approximately 2 zettabytes (2×10^{21} bytes, or roughly 1.6×10^{22} bits). By 2020, this had grown to 64.2 zettabytes — a thirty-fold increase in a single decade. In 2023, global data creation reached 120 zettabytes. In 2024, it hit 149 zettabytes. Projections for 2025 suggest 181 zettabytes, and by 2029 the total is expected to triple from 2025 levels. Approximately 90% of all the data that exists in the world today was generated in the last two years alone. In 2024, approximately 402.74 million terabytes of data were generated every single day — the equivalent of creating the entire Library of Congress every few minutes. Video content accounts for approximately 82% of all global data traffic, driven by streaming services (Netflix, YouTube, TikTok), social media platforms, and user-generated content. Google processes approximately 8.5 billion

searches per day. Over 347 billion emails are sent daily, of which an estimated 70-80% are spam or automated marketing messages. Social media platforms host billions of posts, photos, and videos uploaded each day.

If we convert these figures to Sagan's letter scale, the transformation is remarkable. At 149 zettabytes of data created in 2024 (approximately 1.19×10^{24} bits), Earth in 2024 sits at approximately Type S on Sagan's scale — a jump of eleven full letter categories from the Type H rating Sagan assigned in 1973. In just fifty years, humanity has moved from H to S on the information scale, a rate of advancement that far exceeds anything that has happened on the energy scale over the same period. Earth's Kardashev rating increased from approximately 0.7 to 0.74 — a change of 0.04 — while its Sagan rating increased from H to S — a change of eleven full categories. By this measure, the information revolution has been far more transformative than the energy revolution.

But these figures come with an important caveat that bears repeating. Sagan's scale measures “unique bits of information” — not total data generated. Much of the data humanity generates today is not unique information in any meaningful sense. It consists of copies, duplicates, compressible noise, and redundant transmissions. When a million copies of a best-selling book are printed, the unique information content is the book itself, not the

million copies. When a digital photograph is shared across social media, each share does not represent new information. When a streaming video is watched by a billion viewers, the information content is the video, not the billion views.

This distinction between data and knowledge — between the quantity of bits and the quality of information — is one of the most important limitations of the Sagan scale. A library with a thousand copies of the same book does not contain a thousand times more information. Similarly, a civilization that generates vast quantities of data through automated processes — surveillance cameras, sensor networks, algorithmic trading systems — may not have proportionally more knowledge than a civilization that generates less data but with greater depth and originality.

The challenge of distinguishing data from knowledge becomes more acute as artificial intelligence systems begin generating content. Large language models like GPT-4 can produce text that is grammatically correct, factually plausible, and stylistically consistent — but much of it is derivative, recombinant, or hallucinated. Does AI-generated content count as “unique information” on the Sagan scale? The answer depends on whether the information is novel and true, or merely a rearrangement of existing knowledge. As AI systems become more sophisticated, this distinction will become increasingly difficult to make — and increasingly important for accurate civilizational assessment.

Chapter 5

The Thermodynamics of Information

From Maxwell's Demon to Landauer's Limit

The connection between information and physics was first glimpsed in 1867, when the Scottish physicist James Clerk Maxwell proposed a thought experiment that would trouble thermodynamicists for more than a century. Maxwell, who had already unified electricity and magnetism in his famous equations and developed the kinetic theory of gases, imagined a tiny, intelligent being stationed at a trapdoor between two chambers of gas. This “demon,” as it came to be known, could observe the individual molecules of gas and open or close the trapdoor to allow fast molecules to pass in one direction and slow molecules in the other. Over time, the demon would sort the molecules by speed, creating a temperature difference between the two chambers — heat flowing from cold to hot — without expending any work. This would violate the second law of thermodynamics, which states that the entropy of an isolated system always increases and that heat never flows spontaneously from a cold body to a hot one.

Maxwell's demon was not a serious proposal for a perpetual motion machine. It was a thought experiment designed to probe the limits of thermodynamics and to explore the relationship between macroscopic thermodynamic laws and the microscopic behavior of

individual particles. But it raised a profound question that would take more than a century to fully answer: what prevents such a demon from existing? What physical law, if any, ensures that the second law cannot be violated by an entity with sufficient information about the microstate of a system?

The first serious attempt at a resolution came in 1929, when the Hungarian physicist Leo Szilard — who would later conceive the nuclear chain reaction and co-design the first nuclear reactor with Enrico Fermi — published a paper titled “On the Decrease of Entropy in a Thermodynamic System by the Intervention of Intelligent Beings.” Szilard simplified Maxwell’s thought experiment to a single molecule in a box, divided by a movable partition. The demon determines which side of the partition the molecule is on, then uses the pressure of the molecule pushing against the partition to extract work — $kT \ln 2$ of energy, where k is Boltzmann’s constant and T is the temperature. Szilard argued that the act of measurement — of acquiring information about the molecule’s position — must cost at least as much energy as the work extracted, thereby preserving the second law. Szilard’s analysis was the first to explicitly connect information and energy, showing that information is not merely an abstract mathematical concept but a physical quantity with thermodynamic consequences. Szilard’s analysis was refined over the following decades. In 1932, John von Neumann — the Hungarian-American mathematician

who made foundational contributions to quantum mechanics, computer science, game theory, and virtually every field he touched — proposed an equivalence between information and entropy in his landmark book *Mathematical Foundations of Quantum Mechanics*. Von Neumann showed that the entropy of a quantum system is mathematically identical to the Shannon entropy of an information source, suggesting that the subjective notion of knowledge in an observer’s mind has a deep connection to the objective notion of entropy in physics. Von Neumann’s entropy formula, $S = -\text{Tr}(\rho \ln \rho)$, where ρ is the density matrix of a quantum state, remains the standard definition of entropy in quantum information theory.

In 1948, Claude Elwood Shannon, a mathematician working at Bell Labs, published “A Mathematical Theory of Communication” in the *Bell System Technical Journal* — a two-part paper that founded information theory as a rigorous mathematical discipline. Shannon quantified information in bits (binary digits), defined the concept of information entropy as a measure of uncertainty, established the mathematical framework for measuring and transmitting information, and proved the channel capacity theorem — which sets the theoretical maximum rate at which information can be transmitted over a noisy channel. Shannon’s work is the foundation of all modern digital communication, from the internet to deep space radio links to the compression algorithms that make

streaming video possible. The bit — the unit of information defined by Shannon — is the same unit that Sagan used for his Information Mastery Scale, though Sagan did not explicitly reference Shannon in *The Cosmic Connection*.

In 1961, Rolf Landauer, a physicist at IBM's Thomas J. Watson Research Center in Yorktown Heights, New York, published the paper that would resolve Maxwell's paradox once and for all. Titled "Irreversibility and Heat Generation in the Computing Process" and published in the *IBM Journal of Research and Development*, Landauer's paper established what has come to be known as the Landauer principle: information is physical. It is not an abstract mathematical concept that exists independently of matter and energy. It is encoded in physical systems — the positions of atoms, the orientations of magnetic fields, the states of electronic circuits, the polarization of photons — and it must therefore obey the laws of physics, including the laws of thermodynamics.

From this insight, Landauer derived a fundamental limit on the energy cost of information processing. He showed that while copying information and performing reversible logical operations can, in principle, be done without dissipating energy, there is one operation that must always have a thermodynamic cost: erasing information. When a bit is erased — when two possible states (0 or 1) are collapsed into a single state — the information is lost, and

the entropy of the environment must increase to compensate. The minimum energy required to erase one bit of information is $E = kT \ln 2$, where k is Boltzmann's constant (approximately 1.38×10^{-23} joules per kelvin) and T is the temperature of the environment in kelvin.

At room temperature (approximately 300 kelvin, or 27 degrees Celsius), $kT \ln 2$ works out to approximately 2.85×10^{-21} joules — an extraordinarily small amount of energy, roughly 0.018 electron volts. For comparison, a typical transistor in a modern computer dissipates about 10^{-15} joules per logical operation — approximately one million times the Landauer limit. The energy consumed by a single Google search — approximately 1.5 kilojoules — is enough to erase approximately 5×10^{23} bits at the Landauer limit, more information than is stored in all the world's libraries. This enormous gap between theoretical minimum and practical reality means that modern computers are, from a thermodynamic perspective, extraordinarily inefficient. But it also means that there is vast room for improvement — that computers could, in principle, become millions of times more efficient than they are today.

Landauer's principle was initially met with skepticism. Some physicists argued that it relied on questionable assumptions about statistical mechanics. Others suggested that it might be circumvented by clever engineering that somehow avoids the need

to erase bits. But over the following decades, the principle was repeatedly confirmed experimentally and theoretically.

The first direct experimental verification was published in *Nature* in 2012 by Antoine Berut, Eric Lutz, and colleagues at the Ecole Normale Supérieure in Lyon, France. They used a colloidal particle trapped in a double-well optical potential as a single-bit memory. The particle could be in either the left well (representing 0) or the right well (representing 1). By carefully measuring the heat dissipated when the particle was moved from a random state to a known state — the physical equivalent of erasing a bit — they confirmed that the energy cost matched the Landauer prediction of $kT \ln 2$ within experimental error. The measured heat was approximately $2.8 kT \ln 2$, slightly above the theoretical minimum due to experimental imperfections, but clearly confirming the principle.

Subsequent experiments have confirmed the Landauer principle using different physical systems. In 2016, a team at the University of Maryland used a nanomagnetic bit — a tiny magnetic domain that can be oriented in either of two directions — and measured the heat dissipated during erasure. Other experiments have used electron transport in quantum dots, trapped ions, and superconducting circuits. All have confirmed the basic result: erasing information has a thermodynamic cost that cannot be reduced below $kT \ln 2$ per bit.

In 1973, Charles Bennett — a physicist at IBM’s Thomas J. Watson Research Center and a colleague of Landauer’s — published a landmark paper titled “Logical Reversibility of Computation” in the *IBM Journal of Research and Development*. Bennett showed that reversible computation is theoretically possible: a computer can perform any computation without ever erasing a bit, provided it has enough memory to store intermediate results. In a reversible computer, every logical operation can be undone — the computation can be run backward to recover the input from the output, with no information lost at any step. Such a computer would dissipate far less energy than a conventional computer performing the same computation, because it would avoid the energy cost of bit erasure entirely.

The catch is that reversible computing requires memory — and lots of it. To perform a computation reversibly, the computer must store all intermediate results, because they are needed to run the computation backward. For a complex computation involving many steps, the memory requirements can become enormous. Bennett showed that the memory requirements can be managed by a technique called “pebble games” — periodically saving checkpoints and then reversing the computation back to the last checkpoint — but the fundamental trade-off remains: reversible computing saves energy at the cost of memory.

The implications of Landauer's principle for the Sagan Information Scale are profound. If erasing a bit requires at least $kT \ln 2$ of energy, then processing information at a given rate requires a minimum energy budget, regardless of the technology used. A civilization that processes I bits per second must dissipate at least $I \times kT \ln 2$ watts of heat, assuming all operations are irreversible. This means that there is a fundamental connection between information processing and energy consumption — a connection that the Sagan scale captures implicitly by combining the two dimensions into a single composite rating.

But the connection is not a fixed ratio. A civilization that develops reversible computing can, in principle, process information with arbitrarily small energy cost, limited only by the memory requirements of storing intermediate results. And a civilization that operates at lower temperatures can reduce the Landauer limit proportionally — a computer operating at 3 kelvin (the temperature of the cosmic microwave background in deep space) would have a Landauer limit one hundred times lower than a computer operating at 300 kelvin (room temperature). This is the physical basis for Freeman Dyson's concept of eternal intelligence, which we examine in Chapter 7.

The ultimate limit on computation speed is the Bremermann limit, derived by Hans-Joachim Bremermann in 1962. The Bremermann limit states that the maximum rate at which a system of given mass

and energy can process information is proportional to its energy, divided by Planck's constant: $C_{\text{max}} = m c^2 / h$, where m is the mass of the system, c is the speed of light, and h is Planck's constant. For a one-kilogram computer, this gives approximately 10^{50} operations per second — a figure of the same order of magnitude that Seth Lloyd derived independently in his 2000 paper on the ultimate laptop (Chapter 6). The Bremermann limit is a quantum mechanical limit: it arises from the fact that a quantum system can only change its state at a rate proportional to its energy, per the time-energy uncertainty principle.

Together, the Landauer limit (on energy per operation), the Bremermann limit (on operations per second per kilogram), and the Bekenstein bound (on bits per unit area, discussed in Chapter 7) form the theoretical framework within which any information-processing civilization must operate. Understanding these limits is essential to understanding what the Sagan scale means at the deepest level — and to understanding what is physically possible for civilizations at the highest levels of information mastery.

Chapter 6

The Ultimate Laptop

Seth Lloyd and the Physical Limits of Computation

In August 2000, Seth Lloyd, a professor of mechanical engineering at the Massachusetts Institute of Technology, published a paper in the journal *Nature* that would become one of the most influential works in the physics of computation. Titled “Ultimate Physical Limits to Computation,” the paper asked a deceptively simple question: what are the absolute physical limits on how fast a computer can process information, how much information it can store, and how much energy it requires to do so?

Lloyd’s motivation was the impending end of Moore’s Law. For half a century, the amount of information that computers could process had doubled approximately every eighteen months, driven by the progressive miniaturization of transistors on silicon chips. Gordon Moore, co-founder of Intel, had first observed this trend in 1965, predicting that the number of transistors on an integrated circuit would double every year — a prediction he later revised to every two years, and which held remarkably true for nearly five decades. But this exponential growth could not continue forever. As transistors approached the atomic scale — the most advanced chips as of 2025 feature transistors with critical dimensions of only 3 nanometers, just a few dozen atoms across — quantum effects

begin to dominate, electrons tunnel through barriers that classical physics says they cannot cross, and further miniaturization becomes impossible without fundamentally different technologies. Lloyd's answer was to apply the fundamental constants of physics — the speed of light c , Planck's constant h , Boltzmann's constant k , and the gravitational constant G — to derive absolute limits on computation. These limits, he showed, are set not by engineering constraints but by the laws of physics themselves. No computer, no matter how advanced, can process information faster, store more data, or operate more efficiently than these limits allow. They are the computational equivalent of the speed of light — an absolute ceiling that cannot be exceeded.

To make these limits concrete, Lloyd proposed a thought experiment: the “ultimate laptop.” Imagine a computer with a mass of one kilogram, confined to a volume of one liter — roughly the size and weight of a modern laptop computer. What are the absolute physical limits on what this machine can compute? How many operations per second can it perform? How many bits can it store? How much energy does it consume?

Lloyd's analysis showed that such a machine could process approximately 10^{50} operations per second — a figure so vast that it dwarfs the combined computational power of all the computers ever built on Earth. To put this number in perspective: if you took all 8 billion people on Earth and gave each of them a modern

supercomputer capable of a billion billion (10^{18}) operations per second, the total computing power of all humanity would be approximately 8×10^{27} operations per second. The ultimate laptop could process more than 10^{22} times faster than that — more than ten billion trillion times the combined power of every human and every computer on Earth.

The ultimate laptop could store approximately 10^{31} bits of information — about 10 billion zettabits, or roughly 10 million times the total data generated by humanity in a year. It would dissipate approximately 10^{17} watts of heat during operation — a power output comparable to the total energy consumption of all human civilization (approximately 2×10^{13} watts in 2024). The ultimate laptop, in other words, would be extraordinarily powerful but also extraordinarily hot — it would require cooling systems far beyond anything we can currently imagine.

These numbers are not arbitrary. They are derived directly from fundamental physics. The speed limit of 10^{50} operations per second comes from quantum mechanics: a system of given energy E can change its state at a rate proportional to E / h , where h is Planck's constant. For a one-kilogram system, $E = mc^2 = 9 \times 10^{16}$ joules, giving a maximum rate of approximately 10^{50} operations per second. The storage limit of 10^{31} bits comes from entropy: a system of given size and energy can encode at most $S / k \ln 2$ distinguishable quantum states, where S is the entropy. For a

one-kilogram, one-liter system at maximum entropy, this gives approximately 10^{31} bits. And the energy dissipation of 10^{17} watts comes from thermodynamics: processing 10^{50} bits per second at a temperature of approximately 10^9 kelvin (the temperature at which a one-kilogram black hole would radiate) requires an energy input of approximately 10^{17} watts.

The ultimate laptop is, of course, a theoretical construct. No real computer can approach these limits. The most powerful supercomputers on Earth — machines like Frontier at Oak Ridge National Laboratory (capable of approximately 10^{18} operations per second), Fugaku in Japan, and Aurora at Argonne National Laboratory — operate at speeds that are more than thirty orders of magnitude below the theoretical limit. Even quantum computers, which exploit quantum superposition and entanglement to perform certain calculations exponentially faster than classical computers, are far from the ultimate limit. Google's Sycamore quantum processor, which achieved "quantum supremacy" in 2019 by performing a specific calculation faster than any classical computer could, operates at speeds measured in thousands of operations per second — impressive for a quantum device, but minuscule compared to the 10^{50} ops/sec limit.

But Lloyd's analysis established something profound: there is an absolute ceiling on computational capability, set by the laws of physics, and advanced civilizations will eventually approach it. A

Kardashev Type II civilization, commanding the energy output of a star (approximately 4×10^{26} watts), could in principle build computers that process approximately 10^{70} operations per second — twenty orders of magnitude beyond the one-kilogram laptop. A Kardashev Type III civilization, commanding the energy of a galaxy (approximately 4×10^{37} watts), could build computers processing approximately 10^{81} operations per second. And a civilization that has learned to harvest the energy of the entire observable universe (approximately 10^{50} watts, mostly in the form of dark energy) could approach the ultimate limit of approximately 10^{93} operations per second — the rate at which the universe itself computes.

Lloyd followed up his 2000 paper with a 2001 paper in *Physical Review Letters* titled “Computational Capacity of the Universe,” in which he asked an even more ambitious question: what is the total computational capacity of the entire observable universe?

His answer was derived by treating the universe itself as a giant quantum computer. The observable universe contains approximately 10^{90} particles, mostly photons and neutrinos from the cosmic microwave background, with smaller numbers of electrons, protons, neutrons, and dark matter particles. These particles can, in principle, be used to perform computations, with the total number of operations limited by the total energy of the universe and the total time available since the Big Bang.

Lloyd calculated that the universe has performed approximately 10^{120} operations since the Big Bang — a figure known as the “Lloyd number” — and can store approximately 10^{90} bits of information at any given time. The 10^{120} operations figure comes from multiplying the total energy of the observable universe (approximately 10^{70} joules, mostly in the form of dark energy) by the age of the universe (approximately 4×10^{17} seconds) and dividing by Planck’s constant. The 10^{90} bits storage capacity comes from the total number of particles in the observable universe, each of which can encode a finite number of quantum states.

These numbers represent the absolute upper limits on what any civilization, no matter how advanced, can compute or know using the resources of the observable universe. They are the computational equivalent of the cosmic horizon — the boundary beyond which no information can travel, no computation can reach, and no knowledge can extend.

For the Sagan scale, Lloyd’s analysis sets the ultimate boundary. A Type Z civilization, with 10^{31} bits of information, would be storing approximately 10^{-60} of the total information capacity of the universe — a tiny fraction, but one that still represents an incomprehensibly vast amount of knowledge. A “3 Q” galactic supercivilization, with 10^{26} bits, would be storing approximately 10^{-64} of the universe’s capacity. And our current Type S

civilization, with perhaps 10^{24} bits of stored information, is using approximately 10^{-66} of the universe's capacity — a humbling reminder of how far we have to go, and how much room there is for growth.

Beyond the specific numbers, Lloyd's work introduced a broader principle that has come to be known as "Lloyd's Hypothesis": everything that is worth understanding about a complex system can be understood in terms of how it processes information. This principle — that the information-processing capabilities of a system determine what it can know, what it can do, and what it can become — is the theoretical foundation of the Sagan Information Scale. If Lloyd's Hypothesis is correct, then the information rating of a civilization is not merely a measure of how much data it stores. It is a measure of how well the civilization understands its environment, how effectively it can predict and manipulate the future, and how deeply it can probe the fundamental structure of reality. A civilization with a higher information rating can, by definition, compute more, know more, and do more than a civilization with a lower rating. Information is not just a quantity. It is a capability.

Chapter 7

Dyson's Eternal Intelligence

Information Preservation Across Cosmic Time

In 1979, the same year that Arthur C. Clarke published *The Fountains of Paradise* (Book 3 of this series), the physicist Freeman Dyson published a paper that would become one of the most influential works in cosmological eschatology — the study of the ultimate fate of the universe and the civilizations within it. Titled “Time Without End: Physics and Biology in an Open Universe,” the paper appeared in the journal *Reviews of Modern Physics* and proposed a radical idea: that intelligent life could survive indefinitely in an open, eternally expanding universe, not by commanding ever more energy, but by learning to process information with ever greater efficiency as the universe cooled. Freeman Dyson was one of the most original physicists of the twentieth century. Born in England in 1923, he made fundamental contributions to quantum electrodynamics, solid-state physics, nuclear engineering, and astrophysics. He was a professor at the Institute for Advanced Study in Princeton, New Jersey, for more than sixty years, and his work spanned an extraordinary range — from the Dyson series in quantum field theory to the Dyson sphere concept (Book 1) to his advocacy for nuclear propulsion for space travel. Dyson was known for his willingness to tackle problems

that other physicists considered too speculative or too far from mainstream research, and “Time Without End” was perhaps his most speculative paper — and, in the long run, one of his most influential.

Dyson’s argument in “Time Without End” was elegant and grounded firmly in known physics. He began by considering the fate of the universe. In 1979, it was not yet known whether the universe was open (expanding forever) or closed (destined to recollapse in a Big Crunch). Dyson considered both possibilities but focused on the open case, which he found more interesting because it raised the possibility of infinite survival. In an open universe, the temperature of the cosmic microwave background — the remnant radiation from the Big Bang that fills all of space — drops toward absolute zero as the expansion continues. The current temperature is approximately 2.7 kelvin, but in 10^{14} years it will be approximately 10^{-13} kelvin, and it will continue to fall forever.

Dyson showed that if an intelligent entity could adapt its metabolic rate to match the falling temperature, slowing its thought processes proportionally as the universe cooled, it could perform an infinite number of computations using only a finite amount of energy. The key insight was derived directly from Landauer’s principle: the energy required to process one bit of information is $kT \ln 2$, which approaches zero as the temperature T approaches zero. An

infinitely patient intelligence, willing to think more and more slowly as the universe grew colder, could therefore think forever on a finite energy budget. The total energy required would be the sum of an infinite series: $k \ln 2$ times the sum of T_n for $n = 1$ to infinity, where T_n is the temperature during the n th computational cycle. If the temperature falls sufficiently rapidly — for example, exponentially — this series converges to a finite value, meaning that infinite computation requires only finite energy.

Dyson calculated that a human-level intelligence — one performing approximately 10^{16} operations per second at biological body temperature (approximately 310 kelvin) — could achieve eternal persistence using approximately 6×10^{34} ergs, or 6×10^{27} joules, of energy. To put this in perspective: the Sun's total energy output over its lifetime is approximately 10^{44} joules. The energy required for eternal human-level intelligence is approximately one ten-millionth of the Sun's annual output — a trivially small amount for a civilization capable of stellar engineering. Even a Kardashev Type I civilization, commanding only the energy resources of a single planet, could easily afford the energy budget for eternal intelligence many times over.

Dyson envisioned a specific mechanism for eternal survival: a cycle of activity and hibernation. During active periods, the intelligence would think, compute, and process information at a rate proportional to the ambient temperature. During hibernation

periods, metabolic activity would nearly halt, preserving memory and identity while consuming negligible energy. The duration of active periods and hibernation periods would both scale inversely with temperature — longer and longer periods of dormancy as the universe grew colder — but the subjective experience would be one of continuous consciousness, with no awareness of the intervening eons.

The physical substrate of such an intelligence need not be biological. Dyson argued that advanced civilizations would encode their consciousness into durable forms of matter — perhaps diffuse clouds of particles communicating via electromagnetic signals, or solid-state systems operating at cryogenic temperatures — that could survive the death of stars, the decay of galaxies, and the dilution of matter across the expanding cosmos. The only requirement was thermodynamic reversibility: the computational processes must be reversible, so that no information is lost and no unnecessary energy is dissipated.

The discovery of cosmic acceleration in 1998 — the observation that the expansion of the universe is not slowing down, as gravity would predict, but speeding up, driven by a mysterious force called dark energy — has confirmed that the universe is indeed open and will expand forever. The 2011 Nobel Prize in Physics was awarded to Saul Perlmutter, Brian Schmidt, and Adam Riess for this discovery, which was based on observations of distant Type Ia

supernovae. Dyson's scenario, which was contingent on an open universe, is therefore consistent with our current understanding of cosmology — though the details may need to be modified to account for the effects of dark energy, which may ultimately tear apart galaxies, stars, and even atoms in a “Big Rip” scenario.

Dyson's eternal intelligence connects to the Sagan Information Scale in a profound and direct way. Dyson's intelligence does not command vast energy resources. It does not build Dyson spheres or colonize galaxies. It does not expand spatially or construct megastructures. Instead, it preserves and processes information — the one resource that does not diminish as the universe expands. In Dyson's scenario, the ultimate fate of intelligence is not to become a Kardashev Type III civilization, commanding the energy of a galaxy. It is to become an information-processing entity of extraordinary patience, operating at ever-slower speeds as the universe cools, but maintaining continuous consciousness and memory across infinite time.

This is, in a sense, the logical endpoint of the Sagan scale. If information is the most important resource, and if energy becomes increasingly scarce as the universe expands, then the most advanced civilizations will be those that have learned to maximize information processing per unit of energy — those that have pushed the efficiency of computation toward its theoretical limits. The Sagan scale, read in the light of Dyson's analysis, is not

merely a classification of current capabilities. It is a map of possible futures — a guide to the different ways that civilizations might choose to navigate the long descent toward the heat death of the universe.

A related concept that reinforces the primacy of information is the holographic principle, proposed by the Dutch physicist Gerard 't Hooft in 1993 and developed by Leonard Susskind at Stanford University in 1995, building on earlier work by Jacob Bekenstein on black hole thermodynamics. The holographic principle states that the maximum amount of information that can be contained within a region of space is proportional not to its volume but to the area of its boundary — specifically, one bit of information per Planck area (approximately 10^{-70} square meters).

This means that the information content of any physical system — a black hole, a star, a planet, a computer, a human brain — is bounded not by how much stuff it contains but by the area of its surface. A black hole, paradoxically, encodes all the information about the matter that fell into it on its event horizon — a two-dimensional surface surrounding a three-dimensional volume. This insight, which emerged from the study of black hole thermodynamics by Bekenstein and Stephen Hawking in the 1970s, has led to the suggestion that our entire three-dimensional universe might be a holographic projection of information encoded

on a distant two-dimensional boundary — the “holographic universe” hypothesis.

For the Sagan scale, the holographic principle sets an absolute upper bound on the information that any civilization can store, regardless of how much energy it commands or how much matter it controls. A civilization that has filled a volume of space with information-storage devices cannot exceed the holographic bound for that volume. This bound is so enormous — approximately 10^{70} bits per square meter of surface area — that no conceivable civilization would approach it in practice. But it is a fundamental limit, set by the geometry of spacetime itself, and it reminds us that even information is not infinite — that the universe, vast as it is, has a finite capacity for knowledge.

Chapter 8

Criticisms and Limitations

Does Information Equal Knowledge?

The most obvious criticism of the Sagan Information Scale is that it conflates information with knowledge. Sagan's scale measures the total number of bits a civilization has access to — but not all bits are equally valuable. A civilization that has stored 10^{24} bits of data, mostly consisting of duplicate social media posts, surveillance footage, algorithmically generated content, and redundant backups, does not necessarily know more than a civilization with 10^{18} bits of carefully curated scientific knowledge. The scale counts bits, not insights.

This criticism is particularly relevant to the modern era. The explosion of data generation since 2000 has been driven largely by automated processes: sensor networks, surveillance systems, social media platforms, streaming services, and the Internet of Things that produce data at rates far exceeding any human capacity for comprehension or analysis. In 2024, approximately 402.74 million terabytes of data were generated every day. But the vast majority of this data is ephemeral, redundant, or meaningless. Security cameras record hours of empty hallways. Social media algorithms amplify the same content millions of times. Email servers store

billions of spam messages. Streaming platforms deliver the same movies and songs to billions of viewers and listeners.

Sagan was aware of this distinction, which is why he specified “unique bits of information” rather than total bits. But the concept of “uniqueness” is difficult to define and even more difficult to measure in practice. What counts as unique? Does a copy of a scientific paper count as unique information, or only the original research? Does a translation of a text into another language add new information, or merely repackage old information in a different form? Does raw observational data from a telescope count as information, or only the theories and models derived from it? Does a dataset of human genomes contain new information for each individual, or is most of it redundant given that human genomes are approximately 99.9% identical?

These questions do not have definitive answers, and they highlight a fundamental ambiguity in the Sagan scale. The scale measures something real — the information capacity of a civilization — but it does not distinguish between high-value information (scientific discoveries that reveal new truths about the universe, artistic creations that expand the boundaries of human expression, philosophical insights that transform how we understand ourselves) and low-value information (spam emails, duplicate social media posts, algorithmic trading data, surveillance footage of empty streets). A civilization that generates vast quantities of

data but little genuine knowledge would rate high on the Sagan scale despite being, in a meaningful sense, poorly informed. This problem is compounded by the rise of artificial intelligence. Large language models like GPT-4, Claude, and Llama can produce text that is grammatically correct, factually plausible, and stylistically consistent — but much of it is derivative, recombinant, or outright hallucinated. When an AI system generates a billion words of text, how much of it represents genuine novel information, and how much is merely a rearrangement of existing knowledge? As AI systems become more sophisticated and more deeply integrated into the information ecosystem, this distinction will become increasingly difficult to make — and increasingly important for accurate civilizational assessment on the Sagan scale. A related criticism is cultural bias. The Sagan scale, like the Kardashev and Barrow scales, embeds a particular trajectory of technological development — one that happens to match the history of Western industrial civilization — and presents it as a universal template. Sagan estimated the information content of 1973 Earth by counting the holdings of major libraries — institutions that are predominantly Western in origin and orientation, that collect primarily written and printed materials, and that privilege certain forms of knowledge over others. But other civilizations, past and present, have stored and transmitted knowledge through very different means: oral

traditions, ritual practices, artistic traditions, embodied skills, and non-literate forms of memory that do not translate easily into bits. Aboriginal Australian cultures have maintained songlines — complex oral maps of the landscape encoded in songs and stories — for tens of thousands of years, preserving detailed knowledge of ecology, navigation, water sources, and seasonal patterns across generations without any written records. Polynesian navigators developed sophisticated wayfinding techniques that allowed them to cross thousands of kilometers of open ocean using only their knowledge of star patterns, wave shapes, wind directions, and bird flight paths — a body of knowledge that was largely lost when European contact disrupted traditional practices.

These civilizations might rate low on the Sagan scale (few “bits” in the conventional sense of stored text and data) while possessing knowledge of extraordinary practical value and depth. The Sagan scale captures one dimension of civilizational capability — the ability to store and process large quantities of information in external, countable form — but it does not capture the depth, quality, cultural significance, or practical utility of that information. A civilization with 10^{12} bits of deeply understood, practically applied ecological knowledge might be more “advanced” in meaningful ways than a civilization with 10^{18} bits of shallow, fragmented, algorithmically generated content.

A more fundamental criticism concerns the relationship between information and energy — specifically, the efficiency with which information is processed. As we saw in Chapter 5, the Landauer limit establishes a minimum energy cost for information processing: erasing one bit requires at least $kT \ln 2$ joules of energy. This means that a civilization’s information-processing capacity is ultimately constrained by its energy budget, regardless of how efficient its computing technology becomes. But this constraint also works in the other direction: a civilization with a given energy budget can, by developing more efficient computing technology, process more information per unit of energy.

The ratio of information processed to energy consumed — the “information efficiency” of a civilization — is as important as either quantity alone. A civilization that processes 10^{30} bits per second using 10^{20} watts is, in an important sense, more advanced than a civilization that processes 10^{30} bits per second using 10^{25} watts, because the first civilization has achieved greater thermodynamic efficiency in its computation. This efficiency dimension is not captured by the Sagan scale as originally formulated, which measures only the total information processed, not the energy cost of processing it.

But efficiency is central to the STEM Compression Scale, which we examine in Book 8 of this series. The STEM Compression Scale measures advancement by the efficiency with which a

civilization compresses space, time, energy, and matter into ever-more-productive configurations. A civilization that achieves the same computational capability using less energy has compressed the “E” (energy) dimension of its computational substrate — it has achieved greater efficiency. The synthesis of the Sagan scale (information quantity) with the STEM Compression scale (information efficiency) will be one of the key tasks of the multi-dimensional framework developed in Books 12 and 13.

A final criticism concerns the utility of the Sagan scale for SETI. The Kardashev scale generates testable predictions: a Type II civilization should produce an infrared excess from Dyson sphere waste heat, a Type III civilization should alter the luminosity of its galaxy in detectable ways. These predictions can be tested against astronomical observations, and they have been — the G-HAT infrared survey (Book 1) placed meaningful constraints on the abundance of Type III civilizations in the local universe.

The Sagan scale, by contrast, does not generate obvious technosignatures. A civilization’s information content is not directly observable from interstellar distances. We cannot tell, by looking at a distant star system, whether the civilization there has 10^{15} bits or 10^{25} bits of stored information. The Sagan scale describes capabilities that are invisible to astronomical observation.

However, the Sagan scale may be more relevant to SETI than it first appears. A civilization with a high information rating is likely to be a civilization that communicates — that transmits signals, shares knowledge, and reaches out to other intelligences. As Sagan himself argued, the impulse to communicate may be a natural consequence of high information mastery: a civilization that knows a great deal will want to share what it knows, just as a civilization with abundant energy will want to use that energy for ambitious projects. If this is true, then the Sagan scale may help predict which civilizations are most likely to be detectable — not because their information content produces a direct technosignature, but because their knowledge-seeking nature makes them more likely to transmit radio signals, laser pulses, or other artificial emissions that SETI observers can detect.

Chapter 9

The Knowledge Lens

Preparing for the Twelve-Scale Synthesis

After eight chapters of history, analysis, physics, and criticism, we are in a position to assess the Sagan Information Scale fairly and to determine its place in the larger framework we are building across thirteen books.

The scale gets three things undeniably right.

First, it captures a dimension of civilizational capability that the other scales miss. Kardashev measures energy. Zubrin measures space. Clarke measures tools. Barrow measures inward manipulation of matter. Sagan measures knowledge. These are all real and important dimensions, and no single scale can capture them all. The Sagan scale fills a gap that the other scales leave open — it provides a way to think about civilizations not merely as engines of power or instruments of expansion, but as systems for generating, storing, and transmitting knowledge.

Second, it connects to some of the deepest ideas in modern physics. The thermodynamics of information, the Landauer limit on the energy cost of computation, the holographic principle, the Bremermann limit on computation speed, and the Bekenstein bound on information storage — these are not speculative concepts. They are well-established principles that constrain what

any civilization, anywhere in the universe, can know and compute. The Sagan scale, when read in the light of these principles, becomes a map of physical possibility, not merely a classification scheme. It tells us not just where civilizations are, but where they can go — what is physically possible for beings of given mass, energy, and information at the highest levels of mastery.

Third, it is uniquely relevant to the current moment in human history. We are living through an information revolution — the transition from an energy-limited civilization to an information-saturated one. The exponential growth of data, the rise of artificial intelligence, the explosion of computational power, and the proliferation of global digital networks are the defining developments of our era. The Sagan scale provides a framework for understanding these developments in their full cosmic context. When Sagan proposed his scale in 1973, he estimated Earth at Type H — 10^{13} bits. Half a century later, we are at Type S — 10^{24} bits, a eleven-category jump in a single human lifetime. No other scale in this series shows such dramatic movement for contemporary humanity. By the Sagan measure, we are in the midst of the most rapid transformation in civilizational history. But the scale also has three fundamental limitations that the subsequent books in this series will address.

First, it conflates data with knowledge. The scale measures total bits, not the quality, depth, or usefulness of those bits. A

civilization that generates vast quantities of low-value data — spam, duplicates, noise, surveillance footage of empty streets — would rate high on the Sagan scale despite having limited genuine knowledge. This is not merely a theoretical concern. It is the central challenge of the information age, and it will become increasingly acute as artificial intelligence systems generate ever-larger quantities of content.

Second, it does not measure information efficiency. A civilization that processes 10^{25} bits using 10^{15} watts is more efficient — more advanced, by some measures — than one that processes 10^{25} bits using 10^{20} watts, but the Sagan scale would rate them identically. The efficiency dimension is essential for understanding the long-term sustainability of civilizational development, and it is central to the STEM Compression Scale examined in Book 8.

Third, it does not generate direct technosignatures for SETI observation. Unlike the Kardashev scale, which predicts specific observational signatures, the Sagan scale does not tell astronomers what to look for in the sky. However, high-information civilizations may be more likely to communicate, making them potentially more detectable through their transmissions.

With the completion of this book, we have now examined five of the twelve civilization scales. Let us survey the terrain we have covered.

The Kardashev Scale (Book 1) is the energy lens. It asks: how much power can a civilization command? It connects to the physics of thermodynamics, to the search for megastructures in the sky, and to the fundamental question of whether energy is the right metric for measuring civilizational advancement.

The Zubrin Scale (Book 2) is the spatial lens. It asks: how far can a civilization expand? It connects to the physics of propulsion, to the sociology of settlement, and to the fundamental question of whether growth in spatial reach is the defining feature of advanced civilizations.

The Clarke Scale (Book 3) is the tool lens. It asks: what can a civilization build? It connects to the history of technology, to the magic problem in SETI (Clarke's Third Law), and to the fundamental question of whether tool sophistication captures what makes civilizations advanced.

The Barrow Scale (Book 4) is the inward lens. It asks: how small can a civilization control? It connects to the physics of the very small, to nanotechnology and biotechnology, and to the fundamental question of whether mastery of the microcosm is as important as mastery of the macrocosm.

The Sagan Scale (Book 5) is the knowledge lens. It asks: how much can a civilization know? It connects to the physics of information, to the thermodynamics of computation, and to the

fundamental question of whether knowledge is more significant than power.

Together, these five scales provide a rich, multi-dimensional picture of civilizational capability. A civilization can be rated on all five scales simultaneously: high on Kardashev (lots of energy) but low on Sagan (little knowledge), high on Clarke (sophisticated tools) but low on Barrow (hasn't mastered nanotechnology), high on Zubrin (has colonized widely) but low on Kardashev (uses energy inefficiently), or any other combination. This multi-dimensionality is the core insight that drives the rest of the series, and it will be formalized in the synthesis of Books 12 and 13.

In Book 6, we turn to the Dyson Computation Scale — a framework that measures civilizational advancement not by information stored, but by computation performed. Where Sagan asked “how much do you know?”, Dyson asked “how fast can you think?” The distinction between information and computation is subtle but important: information is static (bits stored), while computation is dynamic (operations performed per second). The Dyson scale, building on the work of Seth Lloyd and the physics of computational limits explored in Chapter 6 of this book, provides a framework for understanding the processing power of civilizations — and for exploring the profound connection between computation, consciousness, and the ultimate fate of intelligence in the universe.

But that exploration is still to come. For now, we have traced the Sagan Information Scale from its origins in a 1973 bestseller, through the physics of information that gives it theoretical depth, to its place in the growing synthesis of civilization classification frameworks. We have seen why information may be, as Sagan believed, “more significant” than energy as a criterion for measuring civilizational advancement. And we have added a fifth lens to our toolkit — a way of seeing civilizations not merely as engines of power, but as engines of knowledge.

Carl Sagan died on December 20, 1996, at the age of sixty-two, after a two-year battle with myelodysplasia, a rare bone marrow disease. He did not live to see the internet revolution transform global communication, the smartphone put a supercomputer in every pocket, social media create new forms of collective intelligence and collective madness, or the artificial intelligence boom that is currently reshaping every aspect of human knowledge work. But the scale he proposed — almost as an aside in a popular book — has proven remarkably prescient. In an age when humanity generates more data in a single day than Sagan could have imagined existing in all the world’s libraries, the question of how to measure what we know — and what knowing means — has never been more urgent. The Sagan Information Scale does not answer that question. But it frames it with the clarity, imagination,

and cosmic perspective that were Carl Sagan's lifelong gifts to science and to humanity.

Epilogue

The Weight of Knowing

We began this book with a question Carl Sagan posed in 1973: might the amount of information a civilization commands be as meaningful a measure of its advancement as the energy it consumes? Over the preceding chapters, we traced that question from its modest origins in *The Cosmic Connection* through the physics of information, the thermodynamic cost of computation, and the absolute limits set by the laws of the universe itself. We watched humanity race from Type H to Type S in barely half a century. We confronted the uncomfortable gap between data and knowledge. And we stood at the edge of what even the most optimistic projections of computation and intelligence can achieve within the observable universe.

What, then, have we learned?

The first lesson is that information is not merely a supplement to energy as a metric of civilization. It is something more fundamental and more slippery. Energy is relatively straightforward to measure. Information, once we move beyond raw bits, becomes entangled with meaning, truth, and usefulness. A civilization can drown in data while starving for knowledge. The Sagan scale, for all its elegance, measures the former more easily than the latter. This is not a flaw in Sagan's original proposal so

much as a reminder that every lens we apply to civilization reveals some things and obscures others.

The second lesson is that the relationship between information and energy is not fixed. Landauer's principle and the work of Seth Lloyd show us that there are hard physical boundaries, but they also reveal that intelligence and computation can, in principle, become dramatically more efficient. A civilization that learns to process information reversibly, or at very low temperatures, or by harnessing the quantum structure of reality itself, could achieve extraordinary informational mastery with relatively modest energy budgets. This possibility undercuts any simple assumption that high information ratings must always be accompanied by high energy ratings. The two axes can diverge.

The third lesson is perhaps the most unsettling. As artificial intelligence systems grow more capable of generating plausible, novel-seeming content, the very category of "unique information" becomes harder to define. If a large language model produces a new scientific hypothesis, a poem, or a philosophical argument that no human has previously expressed, does that count toward a civilization's information rating? The Sagan scale, rooted in a pre-AI understanding of information, does not yet have a clear answer. This is not a problem that can be solved by better counting methods. It is a conceptual problem that forces us to ask what we actually value when we speak of knowledge.

These limitations do not invalidate the Sagan scale. They clarify its proper use. Like the four lenses that preceded it, this one is most powerful when held alongside the others. A civilization that rates high on energy but low on information may be powerful but brittle. A civilization that rates high on information but low on spatial reach may know a great deal yet remain vulnerable to local catastrophes. A civilization that has mastered the very small but processes information inefficiently may achieve technical wonders while remaining conceptually impoverished. Only when we look through multiple lenses at once do we begin to see the full shape of what a civilization is and what it might become.

There is one further implication worth carrying forward. Freeman Dyson's vision of eternal intelligence — civilizations that survive the heat death of the universe not by commanding ever more energy, but by thinking ever more slowly and efficiently as the cosmos cools — is not merely a speculation about the far future. It is a reminder that information processing, at its deepest level, is about the persistence of pattern and meaning against the tide of entropy. In that sense, the Sagan scale points beyond mere classification. It gestures toward something older and more enduring: the ancient human project of preserving what is worth knowing across time.

We are not yet at the end of this series. Six more lenses remain, along with the work of synthesis that will attempt to hold them

together. But with the addition of the knowledge lens, the picture is already richer than any single scale could provide. Energy shows us what a civilization can do. Space shows us where it can go. Tools show us how it acts upon the world. Inward mastery shows us how deeply it can reach into matter. Information shows us what it can remember, understand, and imagine.

The civilizations that will matter most in the long run may not be those that maximize any one of these dimensions. They may be those that learn to balance them — that treat knowledge not as an end in itself, but as a means of becoming more fully what they are capable of becoming, without losing sight of the larger cosmos in which all such becoming takes place.

That, at least, is one way of carrying Sagan's question forward.

Letter to the Reader

Dear Reader,

If you have stayed with this book to the end, you have traveled a considerable distance — from the modest proposal Carl Sagan tucked into a 1973 bestseller, through the thermodynamics of computation, the limits of the observable universe, and the strange new territory opened by artificial intelligence. You have watched humanity’s information rating climb eleven full categories in barely fifty years. You have seen how much of what we now call “information” may not deserve the name.

I wrote this book because the question of what a civilization *knows* — and how that knowledge relates to what it can do — feels increasingly urgent. We are living through an information explosion that Sagan could not have fully anticipated. The tools we now have for generating, storing, and transmitting data are transforming every aspect of human life, often faster than we can make sense of the changes. In such a moment, it seemed worth asking whether one of the twentieth century’s most thoughtful scientists had already given us a useful way to think about what is happening.

I do not claim that the Sagan scale, or any of the scales in this series, offers a complete picture. Each lens we have examined so far leaves important things out of focus. That is precisely why the project continues. The goal has never been to crown a single best

way of measuring civilizations, but to develop a set of complementary tools that, taken together, might help us see more clearly what we are doing and what we might yet become.

If this book has done its job, it has left you with more questions than answers. That is as it should be. The most valuable frameworks are not those that settle every issue, but those that sharpen the questions worth asking. I hope the Sagan scale, in its strengths and in its limitations, has done some of that sharpening for you.

Thank you for reading. The remaining books in the series will continue to add new angles of vision. I hope you will stay with the project as it unfolds.

Sincerest Regards,

James Edmund Carpenter II

Notes on Sources and Further Reading

Primary Sources

Sagan, Carl (1973). *The Cosmic Connection: An Extraterrestrial Perspective*. Anchor Press/Doubleday. The original source of both the decimal Kardashev extension and the Information Mastery Scale (Types A through Z), including the composite “0.7 H” classification for Earth and Sagan’s prediction that first contact would involve civilizations rated “1.5 J” to “1.8 K.”

Key Reviews and Extensions

Gray, Robert H. (2020). “The Extended Kardashev Scale.” *The Astronomical Journal*, 159(5), 228. Contains the formalization of the information scale equation $K_I = (\log_{10}(I) - 6) / 10$, bibliographic investigation tracing the Sagan formula’s origins, and updated estimates of terrestrial information holdings placing Earth at approximately $K_I = 0.8-0.9$ in the 1997-2000 period.

Physics of Information

Shannon, Claude E. (1948). “A Mathematical Theory of Communication.” *Bell System Technical Journal*, 27(3), 379-423 and 27(4), 623-656. Founded information theory; quantified information in bits; established the channel capacity theorem.

Landauer, Rolf (1961). “Irreversibility and Heat Generation in the Computing Process.” *IBM Journal of Research and Development*,

5(3), 183-191. Original paper establishing the Landauer limit: minimum energy $kT \ln 2$ to erase one bit.

Berut, Antoine; Arakelyan, Artak; Petrosyan, Artyom; Ciliberto, Sergio; Dillenschneider, Raoul; and Lutz, Eric (2012).

“Experimental Verification of Landauer’s Principle Linking Information and Thermodynamics.” *Nature*, 483, 187-189. First direct experimental confirmation of the Landauer limit using a colloidal particle in a double-well potential.

Bennett, Charles H. (1973). “Logical Reversibility of Computation.” *IBM Journal of Research and Development*, 17(6), 525-532. Demonstrated that reversible computation is theoretically possible, avoiding the Landauer limit for the computation itself.

Szilard, Leo (1929). “On the Decrease of Entropy in a Thermodynamic System by the Intervention of Intelligent Beings.” *Zeitschrift fuer Physik*, 53, 840-856. Original analysis of Maxwell’s demon and the entropy cost of measurement, establishing $kT \ln 2$ as the fundamental energy scale of information.

Computational Limits

Lloyd, Seth (2000). “Ultimate Physical Limits to Computation.” *Nature*, 406, 1047-1054. Derived the fundamental limits on computation speed (10^{50} ops/sec for 1 kg), storage (10^{31} bits), and energy for the “ultimate laptop.”

Lloyd, Seth (2001). “Computational Capacity of the Universe.” *Physical Review Letters*, 88(23), 237901. Calculated the total computational capacity of the observable universe — approximately 10^{120} operations since the Big Bang (the “Lloyd number”) and 10^{90} bits of storage.

Bremermann, Hans-Joachim (1962). “Optimization through Evolution and Recombination.” In *Self-Organizing Systems*, M.C. Yovits et al. (eds.), Spartan Books. Derived the Bremermann limit on computation speed: $C_{\text{max}} = m c^2 / h$.

Dyson’s Eternal Intelligence

Dyson, Freeman J. (1979). “Time Without End: Physics and Biology in an Open Universe.” *Reviews of Modern Physics*, 51(3), 447-460. Original proposal for eternal intelligence through adaptive metabolism in a cooling universe, requiring approximately 6×10^{27} joules for human-level eternal consciousness.

Holographic Principle

Bekenstein, Jacob D. (1973). “Black Holes and Entropy.” *Physical Review D*, 7(8), 2333-2346. Established the proportionality between black hole entropy and horizon area.

’t Hooft, Gerard (1993). “Dimensional Reduction in Quantum Gravity.” arXiv:gr-qc/9310026. Original proposal of the holographic principle.

Susskind, Leonard (1995). “The World as a Hologram.” *Journal of Mathematical Physics*, 36(11), 6377-6396. Developed the holographic principle in the context of string theory.

Modern Data Growth

Statista (2024). “Data generation volume worldwide 2010-2029.” Reports 149 zettabytes created in 2024, projected 181 zettabytes by 2025.

IDC (2024). Global DataSphere forecast. Projects data generation tripling between 2025 and 2029.

Lesk, Michael (1997). “How Much Information Is There in the World?” Unpublished estimate placing total human knowledge at approximately 10^{18} bits by the late twentieth century.

Rolling Notation Guide (Series-Wide)

This notation guide will be maintained and expanded across all thirteen books in the Civilization Scales Series.

K — Kardashev rating. $K = (\log_{10}(P) - 6) / 10$, where P is power in watts. Introduced in Book 1.

Z — Zubrin rating, measuring spatial reach. Introduced in Book 2.

C_cat — Clarke category (I through VII), measuring tool sophistication. Introduced in Book 3.

B — Barrow rating (I-minus through Omega-minus), measuring inward manipulation scale. Introduced in Book 4.

S_inf — Sagan information rating (A through Z), measuring information content in bits. Introduced in this book (Book 5). Type $A = 10^6$ bits, $H = 10^{13}$ bits (1973 Earth), $S = 10^{24}$ bits (2024 Earth), $Z = 10^{31}$ bits.

K_S — Sagan-Kardashev energy rating. $K_S = (\log_{10}(P) - 6) / 10$.

K-I Composite — Two-dimensional classification combining energy (K_S) and information (S_{inf}), introduced by Sagan in 1973. Notation “ K_I ” means Kardashev rating K with information rating I . Energy and information axes are not strictly interdependent.

Landauer Limit — Minimum energy to erase one bit: $E_{min} = kT \ln 2 \sim 2.85 \times 10^{-21}$ J at room temperature (300 K). Introduced

in Book 5; connects to Book 6 (Dyson Computation Scale) and Book 8 (STEM Compression Scale).

Bremermann Limit — Maximum computation rate per unit mass: $C_{\text{max}} = m c^2 / h \sim 10^{50}$ ops/sec per kilogram. Introduced in Book 5.

Lloyd Number — $\sim 10^{120}$, the total number of operations performed by the observable universe since the Big Bang. Introduced in Book 5.

Bekenstein Bound — Maximum information storage proportional to surface area (approximately 1 bit per Planck area), not volume. Introduced in Book 5.

Transition Danger Zone — $K \sim 0.7$ to 1.0 , the period when a civilization has acquired enough power to threaten its own survival but not enough coordination to manage those risks. Introduced in Book 1.

K-Z Phase Space — Two-dimensional space of Kardashev rating (energy) vs. Zubrin rating (spatial reach). Introduced in Book 2.

K-B Complementarity — Kardashev's outward scale (largest structures) paired with Barrow's inward scale (smallest structures). Introduced in Book 4.

Efficiency Thesis — The argument that the most advanced civilizations achieve the most computation or capability per unit of energy. Connects to Book 8 (STEM Compression Scale).

End of Book 5

The Civilization Scales Series continues with Book 6: *The Dyson
Computation Scale — Thinking as the Sixth Lens*

